

MASTER PROJECT APPENDIX & COMPREHENSIVE REFERENCE COMPENDIUM

Project Title: Sales & Business Performance Analysis

End-to-End Data Cleansing, Relational SQL Modeling, Business Intelligence Dashboarding & Interactive Analytics

Author / Candidate: Gourav
Academic Degree: Bachelor of Technology in Computer Science & Engineering (Data Analytics)
Academic Session: 2025 – 2026
Repository: <https://github.com/Gourav5099/Sales-dashboard>
Live Application Source: `app.py` (Streamlit + Plotly)
Database Staging: `sql/sales_analysis.db` (SQLite 3)
Enterprise BI Workbook: `SPT.pbix` (Power BI Star Schema)

SECTION 1: ACADEMIC & INSTITUTIONAL ALIGNMENT MATRIX

1.1 Vision & Mission Summary

- Institutional Vision:** To emerge as an epicenter of technical excellence, research, and ethically grounded innovation.
- Institutional Mission:** Deliver cutting-edge engineering curriculum, foster multidisciplinary research, and instill leadership and ethical societal commitment.
- Department Vision:** To be globally recognized for excellence in Computer Science, Data Analytics, and Artificial Intelligence education.
- Department Mission:** Cultivate algorithmic and data engineering mastery, experiential laboratory immersion, collaborative industrial projects, and continuous lifelong upskilling.

1.2 Program Outcomes (POs) & Knowledge/Attitude Profiles (WKs)

Program Outcome	Coverage in Project	WK Profile
PO1: Engineering Knowledge	SALES & BUSINESS PERFORMANCE ANALYSIS — SUMMER TRAINING REPORT Applied statistical mathematics, relational algebra, and data normalization.	WK1, WK2

Program Outcome		WK
Coverage in Project		Profile
Department of Computer Science & Engineering Data Analytics Specialization		Candidate: Arav
PO2: Problem Analysis	Diagnosed data anomalies and formulated the <i>Volume vs. Margin Paradox</i> .	WK3, WK4
PO3: Design/Development	Built a relational database schema, Star Schema dimensional model, and Streamlit app.	WK3, WK5
PO4: Conduct Investigations	Performed multi-dimensional exploratory data analysis and discount elasticity queries.	WK4, WK5
PO5: Modern Tool Usage	Leveraged Excel, SQLite, Power BI Desktop, Tableau, Python 3.13, Pandas, Plotly, Streamlit.	WK4
PO6: Engineer & Society	Assessed how operational analytics optimize commercial supply chain logistics.	WK7
PO7: Environment & Sustainability	Aligned retail inventory management with UN Sustainable Development Goals.	WK7
PO8: Ethics	Maintained rigorous data auditing, zero manipulation, and transparent metric reporting.	WK8
PO9: Individual & Team Work	Managed end-to-end data pipeline lifecycle, source version control, and documentation.	WK6
PO10: Communication	Engineered interactive executive dashboards, formal diagrams, and academic reports.	WK5
PO11: Project Management & Finance	Evaluated unit cost economics, gross margin, AOV, and discount leakage.	WK5
PO12: Lifelong Learning	Formulated migration architectures toward modern cloud warehouses (Snowflake, dbt).	WK8

1.3 Course Outcomes (COs)

- **CO1 (Data Cleansing):** Audit raw transactional records, eliminate duplicates, impute missing values, and standardize formatting using Excel and Python.
- **CO2 (Relational SQL):** Design indexed database schemas and author analytical SQL queries utilizing aggregations, CTEs, and window functions.
- **CO3 (Dimensional Modeling & DAX):** Construct Star Schema architectures and develop dynamic business measures in Power BI using DAX.
- **CO4 (Interactive Dashboarding):** Build responsive executive web dashboards in Streamlit and visual stories in Tableau.
- **CO5 (Commercial Synthesis):** Translate multidimensional quantitative findings into data-backed executive recommendations.

1.4 Sustainable Development Goals (SDGs) Alignment

Department of Computer Science & Engineering | Data Analytics Specialization

Candidate: Gourav

- **SDG 8 (Decent Work & Economic Growth - Target 8.2):** Elevates commercial enterprise productivity through technological upgrades and decision support.
 - **SDG 9 (Industry, Innovation & Infrastructure - Target 9.5):** Implements open-source, serverless data infrastructure applicable to small and mid-sized enterprises.
 - **SDG 12 (Responsible Consumption & Production - Target 12.6):** Optimizes stock procurement, mitigates overproduction, and curtails warehouse inventory waste.
-
-

PART A: SOFTWARE AND TOOLS USED

This appendix outlines the complete hardware, operating system, application software, programming runtimes, libraries, and development tools utilized across the design, engineering, execution, and presentation phases of the **Sales & Business Performance Analysis** project.

A.1 Hardware & Operating System Specifications

- **Host Operating System:** Microsoft Windows 11 Home / Professional (64-bit Architecture, x64-based processor).
- **Processor:** Intel Core i5 / i7 (or AMD Ryzen equivalent) with multi-threading support.
- **System Memory (RAM):** 16.0 GB DDR4/DDR5 high-bandwidth memory (supporting in-memory Power BI VertiPaq tabular caching and Streamlit local server execution).
- **Secondary Storage:** High-speed NVMe Solid State Drive (SSD) with ≥ 5.0 GB dedicated storage for repository assets, database staging, and visualization artifacts.

A.2 Core Software Applications & Analytical Platforms

Software / Platform	Vendor / Organization	Version	Operational Function & Implementation Scope
Microsoft Excel	Microsoft Corporation	Office 365 / 2021	Preliminary exploratory data auditing, sanity checking, conditional formatting rules, duplicate identification, and Pivot Table validation.
SQLite Database Engine	SQLite Development Team	3.45.1+	Serverless, zero-configuration relational database staging (<code>sql/sales_analysis.db</code>). Implemented with B-Tree indexes on categorical and temporal fields for ANSI SQL query execution.
Microsoft Power BI Desktop	Microsoft Corporation	2.128.0+	Enterprise dimensional modeling (Star Schema), VertiPaq columnar compression, dynamic DAX calculation engine, and 3-page interactive visual dashboard creation.
Tableau Desktop / Public	Salesforce, Inc.	2024.1+	VizQL visual analytics, multi-dimensional exploratory data analysis, dual-axis margin comparison charts, and visual storytelling.
Visual Studio Code (VS Code)	Microsoft Corporation	1.90.0+	Primary Integrated Development Environment (IDE) for Python scripting, SQL query editing, Git branch management, and markdown report drafting.
Google Chrome / Edge	Google LLC / Microsoft Corp	128.0+	Modern web browser interface for testing local Streamlit reactive apps (<code>localhost:8501</code>) and headless Chromium PDF report generation engine.

A.3 Programming Languages, Runtimes & Dependency Libraries

- **Primary Language & Runtime:** Python 3.13.x (64-bit).

- **Dependency Specification (`requirements.txt`):**
 - `streamlit ≥ 1.30.0` : High performance reactive web framework providing sidebar filters, dynamic KPI metric cards, and responsive session state management.
 - `pandas ≥ 2.0.0` : Fast in-memory tabular data manipulation, vectorized arithmetic, datetime parsing, GroupBy aggregations, and CSV I/O handling.
 - `plotly ≥ 5.18.0` : Interactive JavaScript-backed declarative charting engine rendering interactive tooltips, hover highlights, and zoomable visual layouts.
 - `matplotlib ≥ 3.8.0` : Foundational 2D visualization library used to construct 300 DPI publication-grade static figures (`visuals/generate_charts.py`).
 - `markdown ≥ 3.10.0` & `pygments ≥ 2.21.0` : Text-to-HTML parser and syntax highlighting engine used to automate academic report compilation.

A.4 Version Control, Collaboration & Staging Environment


- **Git:** Version 2.44+ distributed source control system managing branch tracking, semantic commit history, and code lineage.
- **GitHub:** Cloud repository hosting platform (`https://github.com/Gourav5099/Sales-dashboard`) facilitating public portfolio sharing, cloud CI/CD integration, and one-click Streamlit Community Cloud deployment.
- **Streamlit Community Cloud:** Cloud serverless hosting platform deploying `app.py` for permanent global web access.

A.5 Cross-Platform Capability & Comparative Tool Matrix

Operational Feature	Microsoft Excel	SQLite RDBMS	Microsoft Power BI	Tableau Desktop	Python (Streamlit + Plotly)
Data Scalability Limit	1,048,576 rows	Terabytes (disk)	High (In-memory compressed)	High (Hyper extract engine)	High (RAM bound / Chunked)
Analytical Language	Spreadsheet formulas	ANSI SQL	DAX & Power Query M	VizQL / Calculated Fields	Python / Pandas vectorized
Interactivity Level	Medium (Slicers)	None (CLI / Text)	Very High (Cross-filtering)	Very High (Actions/Filters)	High (Reactive Python state)
Automated Scripting	Low (VBA/Macros)	High (SQL scripts)	Medium (Scheduled refresh)	Medium (Tableau Prep)	Full (Automated CLI/APIs)
Licensing / Cost	Paid Commercial	100% Free / Open Source	Free Desktop / Paid Pro	Paid Commercial	100% Free / Open Source

PART B: PROJECT MODULES & ARCHITECTURAL BREAKDOWN

The **Sales & Business Performance Analysis** system is architected as an end-to-end, six-tier modular pipeline. Each module encapsulates distinct business logic, input criteria, transformation rules, and output deliverables:



B.1 Module 1: Synthetic Data Generation & Controlled Anomaly Engine

- Source Component: data/generate_sales_data.py
- Input Parameters: Configuration parameters (Seed: 42, Row count: 3,000, Start date: 2023-01-01, End date: 2024-12-31).

-
- **Functional Operation:** 1. Procedurally generates realistic commercial retail transaction records across 27 distinct product SKUs, 4 categories, and 4 geographic regions. 2. Injects controlled data quality anomalies to mirror real-world point-of-sale defects:
 - Exact duplicate records (10 rows) simulating network re-submission retries.
 - Null/blank fields (10 rows) across Sales, Profit, Discount, Category, and Region.
 - Case mismatches (e.g., `technology`, `north`) and trailing whitespace (e.g., `'Furniture '`).
 - **Output Deliverable:** `data/raw_sales_data.csv` (3,000 rows, 10 columns) and `data/DATA_QUALITY_NOTES.md`.

B.2 Module 2: Data Cleansing, Hygiene & Transformation Engine

- **Source Component:** Excel Audit & Python Data Cleansing Workflow
- **Input:** `data/raw_sales_data.csv` (3,000 rows)
- **Functional Operation:** 1. **Deduplication:** Audits primary key `Order_ID` and drops exactly 10 duplicate rows, yielding 2,990 unique transactions. 2. **Deterministic Imputation:** Rather than inserting artificial column means, missing financial fields are mathematically re-derived:
 - $\text{Sales} = \text{Quantity} * \text{Base Price} * (1 - \text{Discount})$
 - $\text{Profit} = \text{Sales} - (\text{Quantity} * \text{Wholesale Cost})$
 - $\text{Discount} = ((\text{Qty} * \text{Price}) - \text{Sales}) / (\text{Qty} * \text{Price})$3. **Categorical Normalization:** Applies `TRIM()` and `PROPER()` transformations to sanitize categories and regions into uniform Title Case.
- **Output Deliverable:** `sales_data_cleaning.csv` (2,990 clean, fully populated records).

B.3 Module 3: Relational Database Staging & SQL Analytical Execution Engine

- **Source Component:** `sql/sales_analysis.db` and `sql/schema_and_queries.sql`
- **Input:** `sales_data_cleaning.csv`
- **Functional Operation:** 1. **Schema Initialization:** Constructs table `sales_data` with strict type constraints (`INTEGER`, `REAL`, `DATE`, `VARCHAR`). 2. **Indexing:** Creates B-Tree indexes on `Category`, `Region`, `Order_Date`, and `Product` to achieve $O(\log N)$ query performance. 3. **Analytical Query Suite:**
 - *Query 1:* Aggregates high-level enterprise metrics (\$450,978.56 sales, \$157,745.23 profit, 34.98% margin).
 - *Query 2:* Ranks Top 5 revenue-generating products using `ORDER BY Total_Sales DESC LIMIT 5`.
 - *Query 3:* Computes multi-dimensional Category x Region cross-tabulation.
 - *Query 4:* Analyzes merchandise category volume versus profitability.
 - *Query 5:* Evaluates regional revenue and margin distribution.
- **Output Deliverables:** `sql/sales_analysis.db` and `sql/QUERY_RESULTS.md`.

B.4 Module 4: Dimensional Modeling & Power BI Intelligence Engine

- **Source Component:** `SPT.pbix`
- **Input:** Cleaned transactional table staged from SQLite / CSV.

-
- **Functional Operation:** 1. **Star Schema Architecture:** Decouples transactional facts into `Fact_Sales` and temporal dimensional attributes into `Dim_Calendar` (731 dates from 2023 to 2024). 2. **Relationship:** Configures a 1-to-Many single-directional relationship: `Dim_Calendar[DateKey]` -> `Fact_Sales[Order_Date]`. 3. **DAX Measure Library:** Authors explicit DAX measures for financial base values (`[Total Sales]` , `[Total Profit]` , `[Overall Margin %]`), transaction KPIs (`[Total Orders]` , `[AOV]`), and time intelligence (`[Sales YTD]` , `[Sales SPLY]` , `[YoY Growth %]`). 4. **Canvas Design:** Implements a 3-page interactive layout:
 - *Page 1: Executive Overview* (Headline cards, 24-month trend, category donut).
 - *Page 2: Product & Category Deep-Dive* (Ranked bar charts, margin matrix).
 - *Page 3: Regional Commercial Diagnostics* (Geographic breakdown, customer frequency).
 - **Output Deliverables:** `SPT.pbix` and `powerbi/README.md`.

B.5 Module 5: Interactive Web Application & Dashboard Deployment Engine

- **Source Component:** `app.py`
- **Framework:** Streamlit 1.63 + Plotly Graph Objects + Pandas
- **Functional Operation:** 1. **In-Memory Caching:** Leverages `@st.cache_data` to load `sales_data_cleaning.csv` once, ensuring sub-50ms reactive filtering upon user input. 2. **Responsive Sidebar Controls:** Dynamically populates date sliders, region multi-select dropdowns, and category filters. 3. **Visual Layout:**
 - 5 Top KPI metric tiles displaying dynamically recalculated Sales, Profit, Margin, Order Count, and AOV.
 - Monthly sales trajectory line chart with peak annotation.
 - Category revenue share donut chart with center KPI scorecard.
 - Top 5 products horizontal bar chart with margin labels.
 - Regional sales vs. profit clustered column chart.4. **Self-Service Data Export:** Embeds a collapsible raw data table with one-click CSV download capability.
- **Output Deliverable:** `app.py` runnable locally via `streamlit run app.py` and cloud-deployable on Streamlit Community Cloud.

B.6 Module 6: Publication-Grade Static Charting & Automated PDF Engine


- **Source Components:** `visuals/generate_charts.py` and `docs/generate_pdf.py`
- **Functional Operation:** 1. **Static Visual Generation:** Connects to `sql/sales_analysis.db` and utilizes Matplotlib to render four 300 DPI high-resolution PNG charts (`01_sales_trend.png` , `02_top_5_products.png` , `03_category_sales_donut.png` , `04_regional_performance.png`). 2. **Automated PDF Compilation:** Compiles the markdown documentation into styled HTML and invokes headless Google Chrome / Microsoft Edge to produce:
 - Complete formal report (`SUMMER_TRAINING_REPORT.pdf`).
 - Standalone quick-reference compendium (`MASTER_APPENDIX.pdf`).
- **Output Deliverables:** `visuals/*.png` , `SUMMER_TRAINING_REPORT.pdf` and `MASTER_APPENDIX.pdf`

SECTION 2: END-TO-END ANALYTICS LIFECYCLE &

Department of Computer Science & Engineering | Data Analytics Specialization

Candidate: Gourav

PIPELINE



SECTION 3: COMPLETE DATA DICTIONARY & SCHEMA SPECIFICATION

Cleaned Dataset Source: `sales_data_cleaning.csv` (2,990 valid rows, 10 columns)

Database Table: `sales_data` inside `sql/sales_analysis.db`

Column Name	SQL Data Type	Nullable	Primary/Foreign Key	Example Value	Business Definition & Description
<code>Order_ID</code>	<code>TEXT</code>	NO	Primary Key (Unique)	<code>ORD-10024</code>	Unique transaction identifier generated at point-of-sale.
<code>Order_Date</code>	<code>TEXT</code> (ISO-8601)	NO	FK to <code>Dim_Calendar</code>	<code>2023-01-15</code>	Date on which the commercial order was completed (<code>YYYY-MM-DD</code>).
<code>Customer_ID</code>	<code>TEXT</code>	NO	None	<code>CUST-1045</code>	Unique customer account identifier (380 unique accounts).

Column Name	SOL Data Type	Nullable	Primary/Foreign Key	Example Value	Business Definition & Description
Department of Computer Science & Engineering Data Analytics Specialization					
Product	TEXT	NO	None	Standing Desk Converter	Specific merchandise item purchased (27 distinct SKUs).
Category	TEXT	NO	Indexed Dimension	Furniture	Commercial product line (Technology , Furniture , Electronics , Office Supplies).
Region	TEXT	NO	Indexed Dimension	East	Geographic operating territory (East , North , West , South).
Quantity	INTEGER	NO	None	3	Physical count of units ordered in the transaction (\$1 ≤ Q ≤ 5\$).
Sales	REAL	NO	Fact Measure	749.97	Net gross invoice amount billed after discount: \$Q \times P \times (1 - D)\$.
Discount	REAL	NO	Fact Measure	0.15	Promotional discount applied as a decimal rate (\$0.0 ≤ D ≤ 0.40\$).
Profit	REAL	NO	Fact Measure	224.97	Dollar profit retained after wholesale cost: \$\text{Sales} - (Q \times \text{Cost})\$.

SECTION 4: DATA QUALITY DEFECTS AUDIT & CLEANSING RULES

During the Data Preparation phase, six distinct real-world data quality anomalies were audited and cleansed:

Defect Category	Defect Volume	Root Cause	Cleansing & Imputation Protocol Executed	Cleaned Status
Exact Duplicate Rows	10 records	POS double-click / network retry	Purged exact duplicate transactions by unique Order_ID . Ingested 3,000 \$to\$ retained 2,990 clean rows.	Resolved
Missing Sales Values	2 records	POS serialization timeout	Imputed deterministically: \$\text{Sales} = \text{Quantity} \times \text{Base Price} \times (1 - \text{Discount})\$.	Resolved

Defect Category	Defect	Root Cause	Cleansing & Imputation Protocol	Cleaned
	Volume		Executed	Status
Missing Profit Values	2 records	Ledger batch latency	Imputed deterministically: $\text{Profit} = \text{Sales} - (\text{Quantity} \times \text{Wholesale Cost})$.	Resolved
Missing Discount Values	2 records	Promotion code unmapped	Imputed deterministically: $\text{Discount} = \frac{(\text{Qty} \times \text{Price}) - \text{Sales}}{(\text{Qty} \times \text{Price})}$.	Resolved
Missing Category/Region	4 records	Catalog / Territory lookup drop	Imputed Category from Product Master; imputed Region from Customer profile history.	Resolved
Inconsistent Text Casing	15 records	Case mismatch (<code>technology</code> , <code>north</code>)	Standardized via <code>TRIM()</code> and <code>Title Casing</code> (<code>Technology</code> , <code>Furniture</code> , <code>North</code> , etc.).	Resolved

SECTION 5: EXECUTIVE HEADLINE PERFORMANCE SCORECARD

The empirical analysis of the verified 2,990 transactions produced the following enterprise financial benchmarks:

ENTERPRISE COMMERCIAL SCORECARD	
Metric Indicator	Verified Database Result
Total Gross Revenue	\$450,978.56
Total Net Profit	\$157,745.23
Overall Profit Margin	34.98%
Total Completed Transactions	2,990 Orders
Unique Customer Accounts	380 Accounts
Total Physical Units Moved	7,849 Units
Average Order Value (AOV)	\$150.83 per order
Average Profit Per Transaction	\$52.76 per order
Average Basket Size	2.62 units per order
Average Transaction Discount Rate	14.85%

SECTION 6: CORE EMPIRICAL FINDINGS &

Department of Computer Science & Engineering | Data Analytics Specialization

Candidate: Gourav

DIAGNOSTICS

6.1 The Volume vs. Margin Paradox (Category Breakdown)

- **The Volume Drivers:** **Technology** (\\$189,018.44) and **Furniture** (\\$183,965.24) generate **82.70%** of total enterprise sales.
- **The Margin Disconnect:** While Furniture accounts for 40.79% of revenue, it yields the **lowest profit margin in the company (28.24%)**. High wholesale cost bases combined with clearance discounts exceeding 20% severely compress profit margins.
- **The High-Margin Multiplier:** **Office Supplies** accounts for only 6.77% of revenue (\\$30,524.93) but achieves the **highest profit margin (55.26%)**, serving as the prime cross-selling candidate.

Category	Orders	Units Sold	Gross Sales (\$)	Revenue Share	Net Profit (\$)	Profit Share	Profit Margin (%)
Technology	830	2,122	\\$189,018.44	41.91%	\\$66,695.42	42.28%	35.29%
Furniture	622	1,684	\\$183,965.24	40.79%	\\$51,943.84	32.93%	28.24%
Electronics	597	1,631	\\$46,543.94	10.32%	\\$21,917.84	13.89%	47.09%
Office Supplies	934	2,400	\\$30,524.93	6.77%	\\$16,867.62	10.69%	55.26%
Total	2,990	7,849	\\$450,978.56	100.0%	\\$157,745.23	100.0%	34.98%

6.2 Top 5 Revenue-Generating Products

Rank	Product SKU	Category	Units Sold	Gross Revenue (\$)	Net Profit (\$)	Margin (%)	Strategic SKU Role
1	Standing Desk Converter	Furniture	267	\\$62,784.93	\\$16,059.93	25.58%	Top grossing SKU; high discount vulnerability
2	USB-C Docking Station	Technology	291	\\$41,037.18	\\$13,392.18	32.63%	Enterprise workhorse; high corporate demand
3	Ergonomic Office Chair	Furniture	209	\\$38,728.02	\\$9,886.02	25.53%	High ticket anchor; needs discount capping
4	Noise-Canceling Headset	Technology	322	\\$38,009.04	\\$12,893.04	33.92%	High volume, steady healthy margin

Rank	Product SKU	Category	Units Sold	Gross Revenue (\$)	Net Profit (\$)	Margin (%)	Strategic SKU Role
Department of Computer Science & Engineering Data Analytics Specialization Candidate: Gourav							
5	External SSD 1TB	Technology	328	\\$35,834.91	\\$11,234.91	31.35%	High unit velocity consumer staple

6.3 Geographic Sales Territory Distribution

Region	Orders	Units Sold	Gross Sales (\$)	Revenue Share	Net Profit (\$)	Regional Margin (%)
East	826	2,097	\\$121,757.68	27.00%	\\$42,952.11	35.28%
North	790	2,047	\\$116,403.30	25.81%	\\$40,696.00	34.96%
West	698	1,901	\\$108,201.06	24.00%	\\$38,131.99	35.24%
South	672	1,793	\\$104,223.63	23.11%	\\$35,792.74	34.34%

- Regional Takeaway:** Commercial performance is geographically well-balanced. Regional profit margins remain exceptionally stable (34.3%–35.3%), proving that pricing discipline is uniformly maintained across territories.

SECTION 7: PRODUCTION SQL QUERY SUITE & OUTPUT REFERENCE

Database: `sql/sales_analysis.db` | Table: `sales_data`

Query 1: Overall Business KPI Health

```
SELECT
  ROUND(SUM(Sales), 2) AS Total_Sales,
  ROUND(SUM(Profit), 2) AS Total_Profit,
  ROUND((SUM(Profit) / SUM(Sales)) * 100, 2) AS Overall_Margin_Pct,
  COUNT(DISTINCT Order_ID) AS Total_Orders,
  COUNT(DISTINCT Customer_ID) AS Total_Customers,
  SUM(Quantity) AS Total_Units_Sold,
  ROUND(AVG(Sales), 2) AS Avg_Order_Value
FROM sales_data;
```

Output: Sales: \$450,978.56 | Profit: \$157,745.23 | Margin: 34.98% | Orders: 2,990 | Customers: 380 | Units: 7,849 | AOV: \$150.83

Query 2: Top 5 Revenue SKUs

Department of Computer Science & Engineering | Data Analytics Specialization

Candidate: Gourav

```
SELECT
    Product,
    Category,
    SUM(Quantity) AS Units_Sold,
    ROUND(SUM(Sales), 2) AS Total_Sales,
    ROUND(SUM(Profit), 2) AS Total_Profit,
    ROUND((SUM(Profit) / SUM(Sales)) * 100, 2) AS Profit_Margin_Pct
FROM sales_data
GROUP BY Product, Category
ORDER BY Total_Sales DESC
LIMIT 5;
```

Query 3: Multi-Dimensional Category x Region Cross-Tabulation

```
SELECT
    Category,
    Region,
    COUNT(DISTINCT Order_ID) AS Total_Orders,
    SUM(Quantity) AS Units_Sold,
    ROUND(SUM(Sales), 2) AS Total_Sales,
    ROUND(SUM(Profit), 2) AS Total_Profit,
    ROUND((SUM(Profit) / SUM(Sales)) * 100, 2) AS Profit_Margin_Pct
FROM sales_data
GROUP BY Category, Region
ORDER BY Category ASC, Total_Sales DESC;
```

Query 4: Category Summary & Margin Ranking

```
SELECT
    Category,
    COUNT(DISTINCT Order_ID) AS Total_Orders,
    SUM(Quantity) AS Units_Sold,
    ROUND(SUM(Sales), 2) AS Total_Sales,
    ROUND(SUM(Profit), 2) AS Total_Profit,
    ROUND((SUM(Profit) / SUM(Sales)) * 100, 2) AS Profit_Margin_Pct
FROM sales_data
GROUP BY Category
ORDER BY Total_Sales DESC;
```

Query 5: Regional Territory Summary

Department of Computer Science & Engineering | Data Analytics Specialization

Candidate: Gourav

```
SELECT
    Region,
    COUNT(DISTINCT Order_ID) AS Total_Orders,
    SUM(Quantity) AS Units_Sold,
    ROUND(SUM(Sales), 2) AS Total_Sales,
    ROUND(SUM(Profit), 2) AS Total_Profit,
    ROUND((SUM(Profit) / SUM(Sales)) * 100, 2) AS Profit_Margin_Pct
FROM sales_data
GROUP BY Region
ORDER BY Total_Sales DESC;
```

SECTION 8: POWER BI STAR SCHEMA & DAX MEASURE LIBRARY

8.1 Star Schema Data Architecture

- Fact Table (Fact_Sales)**: 2,990 rows containing transaction keys (Order_ID , Order_Date , Customer_ID , Product , Category , Region) and quantitative metrics (Quantity , Sales , Discount , Profit).
- Dimension Table (Dim_Calendar)**: Continuous date table generated spanning 2023-01-01 to 2024-12-31 (731 rows) with attributes: Year , Quarter , Month , MonthName , DayOfWeek , IsWeekend .
- Relationship**: Fact_Sales[Order_Date] $\rightarrow$ Dim_Calendar[DateKey] (Single directional cross-filtering).

8.2 Production DAX Measure Library

Core Financial Measures

Total Sales = SUM(Fact_Sales[Sales])

Total Profit = SUM(Fact_Sales[Profit])

Overall Profit Margin % = DIVIDE([Total Profit], [Total Sales], 0) * 100

Total Orders = DISTINCTCOUNT(Fact_Sales[Order_ID])

Total Units Sold = SUM(Fact_Sales[Quantity])

SALES & BUSINESS PERFORMANCE ANALYSIS — SUMMER TRAINING REPORT

Average Order Value (AOV) = DIVIDE([Total Sales], [Total Orders], 0)

Department of Computer Science & Engineering | Data Analytics Specialization

Candidate: Gourav

Time-Intelligence & Growth Measures

Sales YTD = TOTALYTD([Total Sales], 'Dim_Calendar'[DateKey])

Sales SPLY = CALCULATE([Total Sales], SAMEPERIODLASTYEAR('Dim_Calendar'[DateKey]))

```
YoY Sales Growth % =  
VAR PreviousYear = [Sales SPLY]  
VAR CurrentSales = [Total Sales]  
RETURN  
    IF(  
        ISBLANK(PreviousYear),  
        BLANK(),  
        DIVIDE(CurrentSales - PreviousYear, PreviousYear, 0) * 100  
    )
```

SECTION 9: INTERACTIVE WEB APPLICATION ARCHITECTURE (`app.py`)

- **Framework:** Streamlit 1.63 + Plotly Graph Objects + Pandas
- **Execution Script:** `streamlit run app.py`
- **Core Architecture Components:**
 1. **Caching Engine:** `@st.cache_data` caches the 2,990-row dataset in memory for instant sub-second filtering.
 2. **Interactive Controls:**
 - Temporal Slider: Dynamic Date Range filtering across 2023–2024.
 - Multi-Select Dropdowns: Filter by `Region` and `Category`.
 3. **Visual Layout:**
 - Top KPI Tiles: Live recalculated Gross Revenue, Net Profit, Profit Margin %, Total Orders, and AOV.
 - Chart 1: 24-Month Time-Series Trend Line (Sales vs. Profit with peak annotations).
 - Chart 2: Merchandise Category Revenue Share Donut Chart with center metric scorecard.
 - Chart 3: Top 5 Revenue SKUs Horizontal Ranked Bar Chart.
 - Chart 4: Clustered Regional Bar Chart comparing Sales vs. Net Profit.
 4. **Data Export:** Collapsible raw data table with live CSV download button for downstream reporting.

SECTION 10: ACTIONABLE STRATEGIC

Department of Computer Science & Engineering | Data Analytics Specialization

Candidate: Gourav

RECOMMENDATIONS

- Implement a 15% Maximum Discount Ceiling on Furniture:** - High-cost items like *Standing Desk Converters* and *Ergonomic Chairs* currently experience severe margin collapse when discounted >20%. Capping discounts at 15% preserves an estimated **\\$4,200 to \\$6,500 in quarterly profit**.
- Checkout Bundling with High-Margin Office Supplies:** - Leverage the **55.26% profit margin** of Office Supplies by automatically offering add-on bundles (desk organizers, cable management sleeves, premium paper) at point-of-sale when hardware is purchased. This expands AOV without diluting hardware margins.
- Targeted Regional B2B Expansion in Southern Territory:** - The Southern region currently lags behind East by \\$17.5K in gross sales despite having identical profit margin health (34.34%). Directing corporate sales reps to Southern accounts captures high-margin untapped demand.
- Tiered B2B Corporate Loyalty Program:** - With 380 unique accounts placing 2,990 orders (~7.8 orders per customer), high customer retention is proven. A tiered rebate incentive based on annual volume protects top accounts against competitor poaching.

SECTION 11: VIVA VOCE DEFENSE & EXAMINER Q&A CHEAT SHEET

#	Typical Examiner Viva Question	Senior Data Analyst Bulletproof Defense Answer
Q1	Why did you choose SQLite over MySQL or PostgreSQL?	SQLite is a serverless, zero-configuration in-process database that reads directly from <code>sales_analysis.db</code> . It provides full ACID compliance, B-Tree indexing, and ANSI SQL window functions with zero server daemon overhead, making the project 100% portable for any examiner without credential setups.
Q2	What is the most critical commercial finding of your project?	The Volume vs. Margin Paradox : Technology and Furniture drive 82.70% of gross revenue, but Furniture has the lowest margin (28.24%) due to high wholesale costs and discount leakage (>20%), whereas Office Supplies generates only 6.77% of sales but delivers a 55.26% margin.
Q3	Why did you create a Star Schema instead of using a single flat table in Power BI?	A Star Schema separates numerical facts (<code>Fact_Sales</code>) from temporal attributes (<code>Dim_Calendar</code>). This optimizes the VertiPaq columnar compression engine, prevents grain mismatch errors, and enables standard Time Intelligence DAX functions (<code>TOTALYTD</code> , <code>SAMEPERIODLASTYEAR</code>).
Q4	How did you handle data hygiene and missing values?	Exactly 10 duplicate rows were identified and purged via unique <code>Order_ID</code> . Missing sales, profit, and discount values were not guessed—they were imputed deterministically using fundamental pricing logic: $\text{Sales} = Q \times P$, $\text{Profit} = \text{Sales} \times (1 - D) - \text{Cost}$, and $\text{Discount} = \text{Sales} - \text{Revenue}$. SALES ANALYSIS REPORT DATA SOURCES: ERP, CRM, POS PREPARED BY: DATA ANALYST DATE: 2024-10-27

#	Typical Examiner Viva Question	Senior Data Analyst Bulletproof Defense Answer
Department of Computer Science & Engineering Data Analytics Specialization		Candidate: Gourav
Q5	What is the difference between an implicit and explicit DAX measure?	Implicit measures are auto-generated by dragging fields into visuals (e.g., auto-sum). Explicit measures are user-defined formulas using DAX functions (e.g., <code>[Total Sales] = SUM(...)</code>). Explicit measures are mandatory for time intelligence, dynamic filtering, and clean visual architecture.
Q6	What does <code>@st.cache_data</code> do in your Streamlit application?	Streamlit re-runs the entire script from top to bottom on every user interaction. <code>@st.cache_data</code> stores the loaded CSV in memory so subsequent slider/dropdown clicks execute in milliseconds without re-reading the disk.
Q7	Why is Office Supplies margin so much higher than Furniture?	Office supplies have low wholesale cost and high consumer markup elasticity. People rarely negotiate discounts on small items like pens or organizers, preserving a 55.26% margin. Furniture items carry heavy manufacturing wholesale costs that collapse when discounted.
Q8	How does your project align with United Nations SDGs?	It directly advances SDG 8 (Decent Work & Economic Growth - Target 8.2) by boosting productivity, SDG 9 (Industry & Innovation - Target 9.5) through low-cost analytics infrastructure, and SDG 12 (Responsible Consumption - Target 12.6) by mitigating inventory overproduction.
Q9	How would you scale this architecture to 50 million rows?	I would migrate storage from SQLite to a cloud data warehouse like Snowflake or Google BigQuery , orchestrate transformation models using dbt , convert CSV storage to columnar Parquet , and connect Power BI via DirectQuery / Aggregation Tables.
Q10	What are the primary limitations of your dataset?	The dataset lacks customer demographic details (age, income bracket), web telemetry (page clicks, bounce rates), and marketing channel attribution (ad spend per order). Adding marketing spend would enable Customer Acquisition Cost (CAC) and ROAS analysis.

SECTION 12: PROJECT FILE REPRODUCIBILITY & ARTIFACT MANIFEST

File / Folder Path	Type	Purpose & Contents
<code>data/raw_sales_data.csv</code>	Data	3,000 raw transactions with controlled quality defects.
<code>data/DATA_QUALITY_NOTES.md</code>	Doc	Detailed defect log of all 10 duplicates, 10 nulls, and formatting issues.
<code>sales_data_cleaning.csv</code>	Data	2,990 clean, standardized, verified transaction records.
<code>sql/sales_analysis.db</code>	Database	SQLite 3 relational database containing indexed <code>sales_data</code> .
<code>sql/schema_and_queries.sql</code>	SQL	SALES & BUSINESS PERFORMANCE ANALYSIS — SUMMER TRAINING REPORT Full DDL table definitions and 5 analytical queries.

File / Folder Path	Type	Purpose & Contents
Department of Computer Science & Engineering Data Analytics Specialization Candidate: Gourav		
SPT.pbix	Power BI	Star Schema dimensional data model and DAX measures.
app.py	Python	Interactive Streamlit + Plotly web dashboard application.
requirements.txt	Config	Dependencies pinned for deployment (streamlit , pandas , plotly , matplotlib).
visuals/generate_charts.py	Python	Standalone script generating 300 DPI publication-grade PNG charts.
visuals/*.png	Visuals	4 high-res charts (Sales Trend, Top Products, Donut, Regional Bar).
documentation/SUMMER_TRAINING_REPORT.md	Doc	Full 92-page academic summer training report.
SUMMER_TRAINING_REPORT.pdf	PDF	Formatted, publication-grade academic PDF report.
documentation/MASTER_APPENDIX.md	Doc	This file: Comprehensive Master Appendix & Executive Compendium.

Terminal Execution Cheat Sheet

```
# 1. Run Interactive Streamlit Web App
streamlit run app.py

# 2. Regenerate High-Resolution Charts
python visuals/generate_charts.py

# 3. Regenerate Full Academic PDF Report
python docs/generate_pdf.py

# 4. Git Synchronization
git add .
git commit -m "Update project documentation"
git push origin main
```